

# Machine Intelligence:: Deep Learning

## Week 2

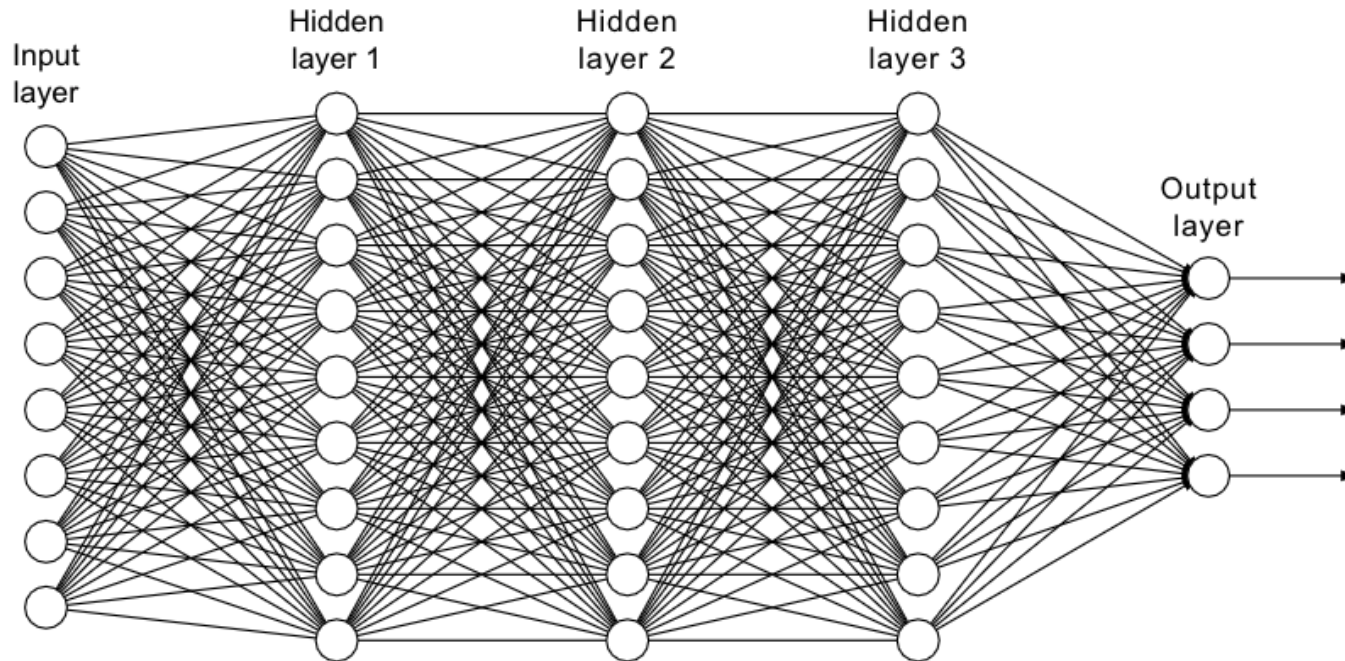
*Beate Sick, Oliver Dürr*

Institut für Datenanalyse und Prozessdesign  
Zürcher Hochschule für Angewandte Wissenschaften

# Topics of today


- A second look on fully connected Neural Networks (fcNN)
- Convolutional Neural Networks (CNN) for images
  - Motivation for switching from fcNN to CNNs
  - Introduction of convolution
  - ReLu and Maxpooling Layer
  - Biological inspiration of CNNs
  - Building CNNs

# Architecture of a fully connected NN




Each neuron in a fcNN gets as input a weighted sum of all neuron activation from one layer below. Different neurons in the same layer have different weights in this weighted sum, which are learned during training.

# Comon non-linear activation function




$$y = f(z) = f(b + \sum x_i \cdot w_i)$$



The sigmoid has small gradients for values far away from zero.  
ReLU clips values below zero and let values  $> 0$  pass unchanged.

# Best practice: Split in Train, Validation, and Test Set




Best practice: Lock an extra **test data set** away, and use it only at the very end, to evaluate the chosen model, that performed best on your validation set.

Reason: **When trying many models, you probably overfit on the validation set.**

Determine performance metrics, such as MSE, to evaluate the predictions **on new validation or test data**

# What can loss curves tell us?

Very common check: Plot loss in train and validation data vs epoch of training.




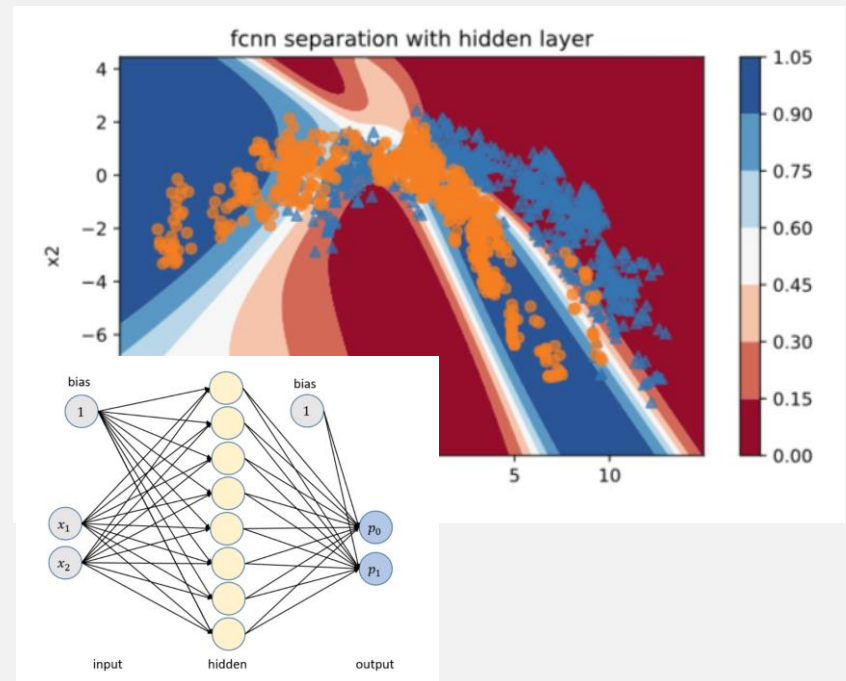

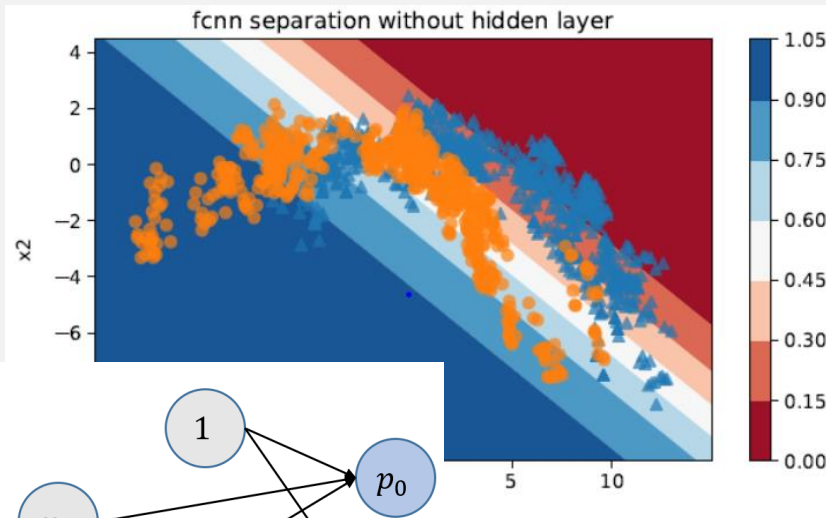

- If training loss does not go down to zero: model is not flexible enough
- In case of overfitting (validation loss  $\gg$  train loss): regularize model

# Homework

Do exercise NB 02: Classify

banknotes based on 2 features ( $x_1, x_2$ )


[https://github.com/tensorchiefs/dl\\_course\\_2022/blob/master/notebooks/02\\_fcnn\\_with\\_banknote.ipynb](https://github.com/tensorchiefs/dl_course_2022/blob/master/notebooks/02_fcnn_with_banknote.ipynb)




Fully connected NN for image data  
Why not?



# A fcNN for MNIST data




The first four digits of the MNIST data set - each image consisting of  $28 \times 28 = 784$  pixels



A fully connected NN with 2 hidden layers.

For the MNIST example, the input layer has 784 values for the  $28 \times 28$  pixels and the output layer has 10 nodes for the 10 classes.

# What is going on in a 1 layer fully connected NN?




Cost C or Loss = cross-entropy averaged over all images in mini-batch


$$C = \frac{1}{N} \sum_i D(\mathbf{p}_i, \mathbf{y}_i)$$

Cross-Entropy

$$D(\mathbf{p}, \mathbf{y}) = - \sum_{k=1}^{10} y_k \cdot \log(p_k)$$



# Exercise: Does shuffling disturb a fcNN?



Use fcNN for MNIST:

[https://github.com/tensorchiefs/dl\\_course\\_2022/blob/master/notebooks/03\\_fcnn\\_mnist.ipynb](https://github.com/tensorchiefs/dl_course_2022/blob/master/notebooks/03_fcnn_mnist.ipynb)

Investigate if shuffling disturbs the fcNN for MNIST:


[https://github.com/tensorchiefs/dl\\_course\\_2022/blob/master/notebooks/04\\_fcnn\\_mnist\\_shuffled.ipynb](https://github.com/tensorchiefs/dl_course_2022/blob/master/notebooks/04_fcnn_mnist_shuffled.ipynb)

# Convolutional Neural Networks


## SoA for image data

# Convolution extracts local information using few weights

Fully connected neural net




Locally connected neural net



## Shared weights:

by using the **same weights for each patch** of the image we need much **less parameters** than in the fully connected NN and get from each patch the same kind of **local feature information** such as the presence of an edge.

# An artificial neuron





# Convolutional networks use neighborhood information and replicated local feature extraction

In a locally connected network the calculation rule

$$z = b + \sum_i x_i w_i$$



Pixel values in a small image patch are element-wise multiplied with weights of a small filter/kernel:




The filter is applied at each position of the image and it can be shown that the **result is maximal if the image pattern corresponds to the weight pattern.**

The results form again an image called **feature map** (=activation map) which shows at which position the feature is present.

feature/activation map




# Applying the same 3x3 kernel at each image position



Applying the 3x3 kernel on a certain position of the image yields one pixel within the activation map where the position corresponds to the center of the image patch on which the kernel is applied.



# Convolutional networks use neighborhood information and replicated local feature extraction



The weights of each filter are randomly initiated and then adapted during the training.


## Exercise: Do one convolution step by hand



The kernel is 3x3 and is applied at each valid position  
– how large is the resulting activation map?

The small numbers in the shaded region are the kernel weights.  
Determine the position and the value within the resulting activation map.

3	3	2	1	0
$0_0$	$0_1$	$1_2$	3	1
$3_2$	$1_2$	$2_0$	2	3
$2_0$	$0_1$	$0_2$	2	2
2	0	0	0	1




# Example of designed Kernel / Filter




$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

Horizontal Sobel kernel




Applying a vertical edge detector kernel

# CNN Ingredient I: Convolution




Zero-padding to achieve  
**same** size of feature and input



no padding to only use  
**valid** input information


The *same* weights are used at each position of the input image.

# Building a very simple CNN with keras




```
model = Sequential()
model.add(Convolution2D(1, (5,5), # one 5x5 kernel
                        padding='same', # zero-padding to preserve size
                        input_shape=(pixel,pixel,1)))
model.add(Activation('linear'))
# take the max over all values in the activation map
model.add(MaxPooling2D(pool_size=(pixel,pixel)))
model.add(Flatten())
model.add(Dense(2))
model.add(Activation('softmax'))
```

# Exercise: Artstyle Lover



Open NB in: [https://github.com/tensorchiefs/dl\\_course\\_2022/blob/master/notebooks/05\\_cnn\\_edge\\_lover.ipynb](https://github.com/tensorchiefs/dl_course_2022/blob/master/notebooks/05_cnn_edge_lover.ipynb)


## Convolution layer with a 1-channel input and 6 kernels



Convolution of the input image with 6 different kernels results in 6 activation maps.  
If the input image has only one channel, then each kernel has also only one channel.


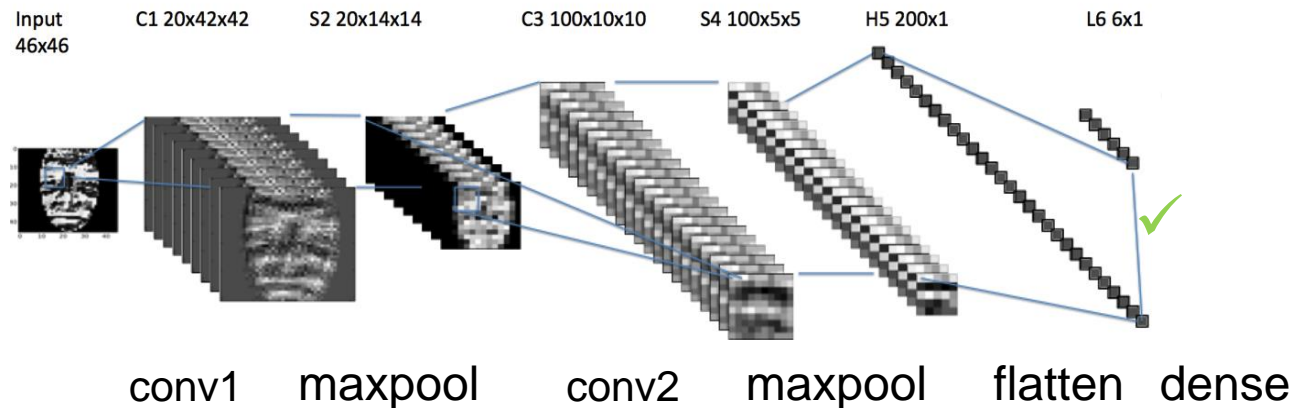


# A CNN with 3 convolution layers






# CNN ingredient II: Maxpooling Building Blocks reduce size



Simply join e.g. 2x2 adjacent pixels in one by taking the max.  
→ less weights in model  
→ Less train data needed  
→ increased performance

Hinton: „The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster“

# Animated convolution with 3 input channels




3 color channel input image

$W[4, 4, 3]$

The value of neuron  $j$  in the  $k$ -th featuremap are computed from the weights in the  $k$ -th filter  $w_{ki}$  and the input values  $x_{ji}$  at the position  $j$ :


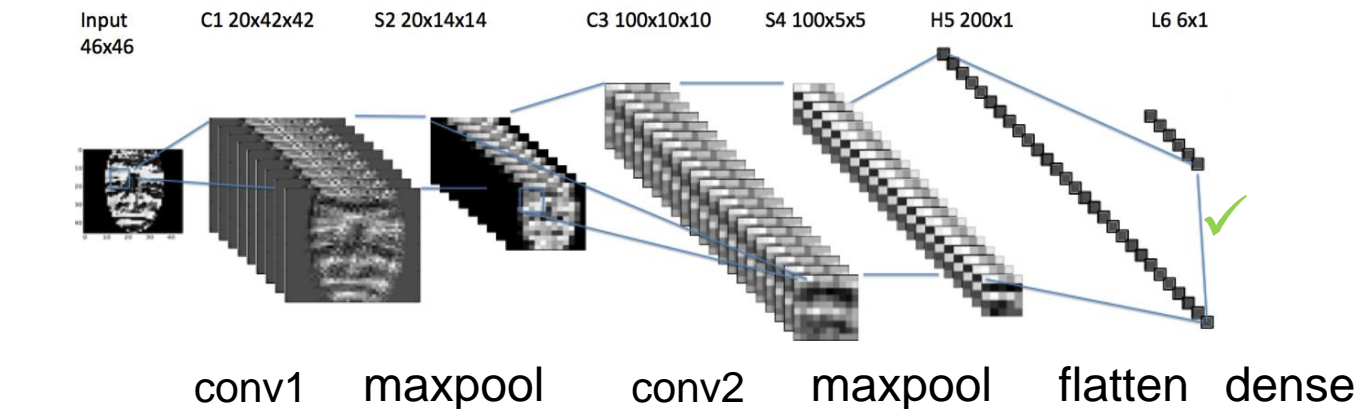
$$y_{jF_k} = f(z_{jF_k}) = f(b_k + \sum x_{ji} \cdot w_{ki})$$

## Convolution layer with a 3-channel input and 6 kernels



Convolution of the input image with 6 different kernels results in 6 activation maps. If the input image has 3 channels, then each filter has also 3 channels.


# CNN ingredient II: Maxpooling Building Blocks reduce size



Simply join e.g. 2x2 adjacent pixels in one by taking the max.  
→ less weights in model  
→ Less train data needed  
→ increased performance

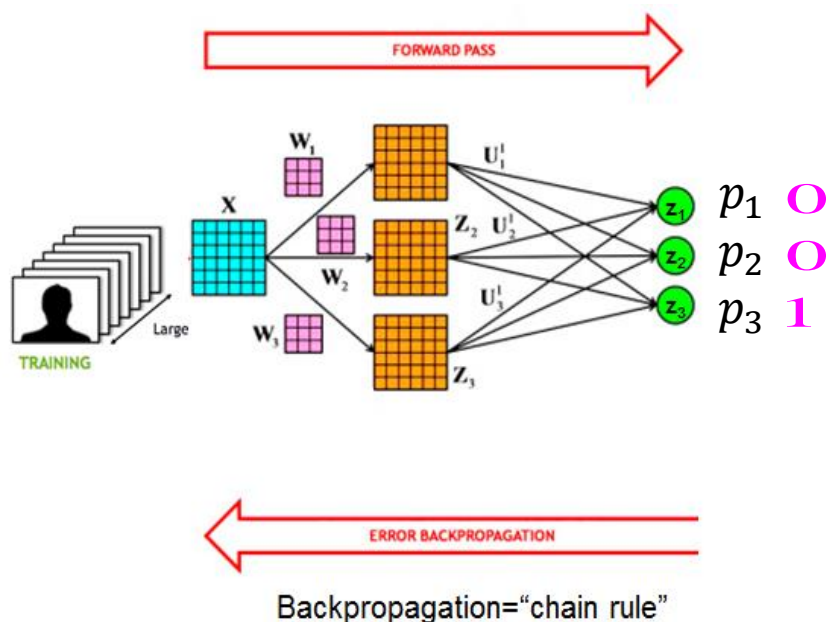
Hinton: „The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster“

# Typical shape of a classical CNN



Spatial resolution is decreased e.g. via max-pooling while more abstract image features are detected in deeper layers.

# Training of a CNN is based on gradient backpropagation



For the training we need the **observed label** for each image which we then compare with the **output** of the CNN.

We want to **adjust the weights** in a way so that difference between true label and output is minimal.


Minimize Loss-function:

**$L = \text{distance}(\text{observed}, \text{output}(w))$**

Learning is done by weight updating:

$$w_i^{(t)} = w_i^{(t-1)} - \underset{\substack{\uparrow \\ \text{learning rate}}}{l^{(t)}} \frac{\partial L(w)}{\partial w_i} \bigg|_{w_i = w_i^{(t-1)}}$$


# CNN for MNIST



```
unclassified = 10          # 10 classes: 0, 1,...,9
input_shape = (28, 28, 1) # 28x28 pixels, 1 channel (grey value)

model = Sequential()
model.add(Convolution2D(32, (3, 3), #32 filters of size 3x3
                        activation='relu',
                        input_shape=input_shape))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Convolution2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(40, activation='relu'))
model.add(Dense(num_classes, activation='softmax'))
```

# Exercise: Does shuffling disturb a CNN?



Open NB in: [https://github.com/tensorchiefs/dl\\_course\\_2022/blob/master/notebooks/06\\_cnn\\_mnist\\_shuffled.ipynb](https://github.com/tensorchiefs/dl_course_2022/blob/master/notebooks/06_cnn_mnist_shuffled.ipynb)




# fcNN versus CNNs – some aspects


- A fcNN is good for tabular data, CNNs are good for ordered data (eg images).
- In a fcNN the order of the input does not matter, in CNN shuffling matters.
- A fcNN has no model bias, a CNN has the model bias that neighborhood matters.
- A node in one layer of a fcNN corresponds to one feature map in a convolution layer:
- In each layer of a fcNN connecting  $p$  to  $q$  nodes, we learn  $q$  linear combinations of the incoming  $p$  signals, in each layer of a CNN connecting  $p$  channels with  $q$  channels we learn  $q$  filters (each having  $p$  channels) yielding  $q$  feature maps

# Biological Inspiration of CNNs


# How does the brain respond to visually recieved stimuli?




Setup of the experiment of Hubel and Wiesel in late 1950s in which they discovered **neurons** in the visual cortex that **responded** when moving **edges** were shown to the cat.



Receptive fields size




Features





Organization of the visual cortex in a brain, where neurons in different regions respond to more and more complex stimuli

# Compare neurons in brain region V1 in first layer of a CNN




V1 physiology: orientation selectivity



Neurons in brain region V1 and neurons in 1. layer of a CNN respond to similar patterns

# Visualize the weights used in filters

Filter weights from a trained Alex Net





Only in layer 1 the filter pattern correspond to extracted patterns in the image.

In higher layers we can only check if patterns look noisy, which would indicate that the network that hasn't been trained for long enough, or possibly with a too low regularization strength that may have led to overfitting.

# The receptive field

For each pixel of a feature map we can determine the connected area in the input image – this area in the input image is called receptive field.


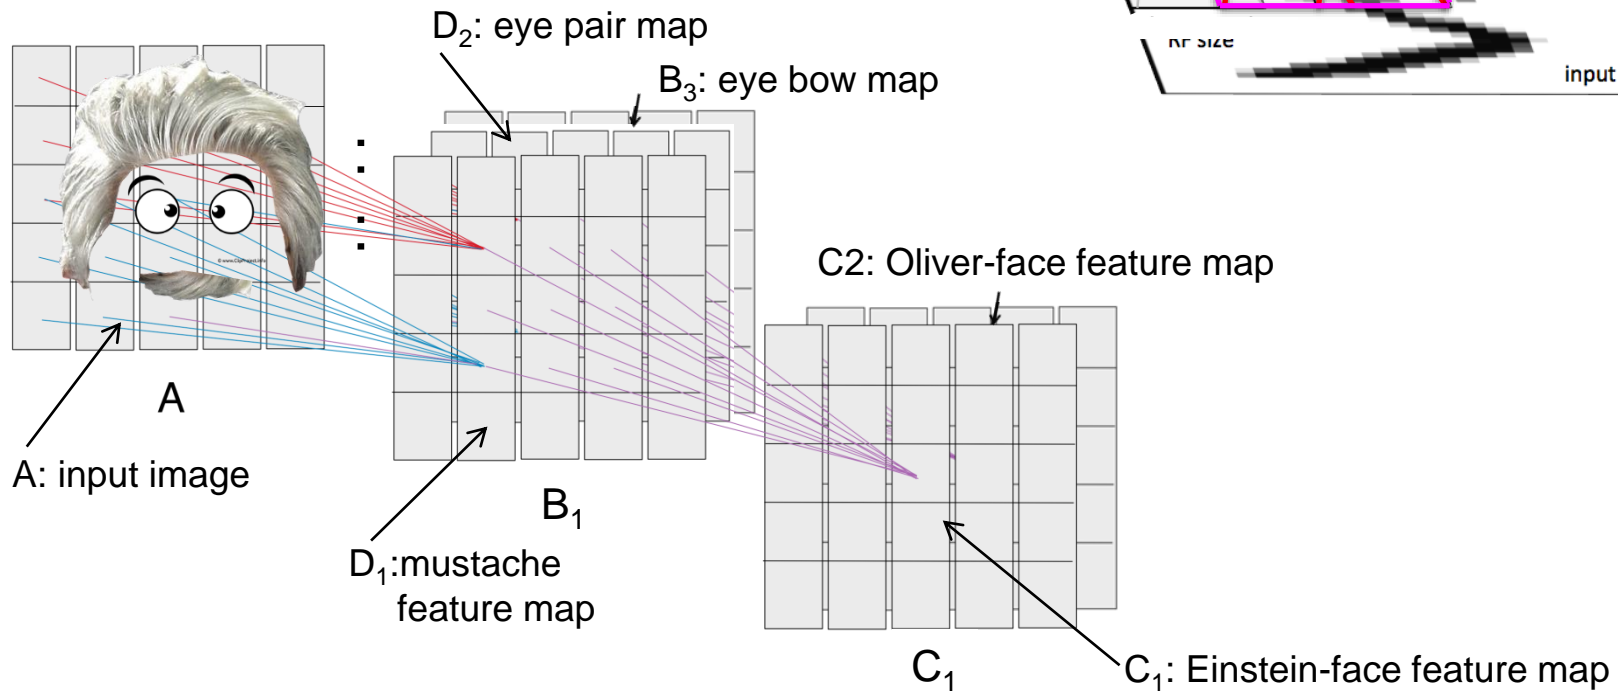


A feature map gets activated by a certain structure of the feature maps one layer below, which by itself depends on the input of a preceding layer etc and finally on the input image. Activation maps close to the input image are activated by simple structures in the image, higher maps by more complex image structures.




# The receptive field

The receptive field gets larger and larger when going further away from the input



Filter cascade across different channels can capture relative position of different features in input image. Einstein-face-filter will have a high value at expected mustache position.

# Weak analogies between brain and CNNs architecture

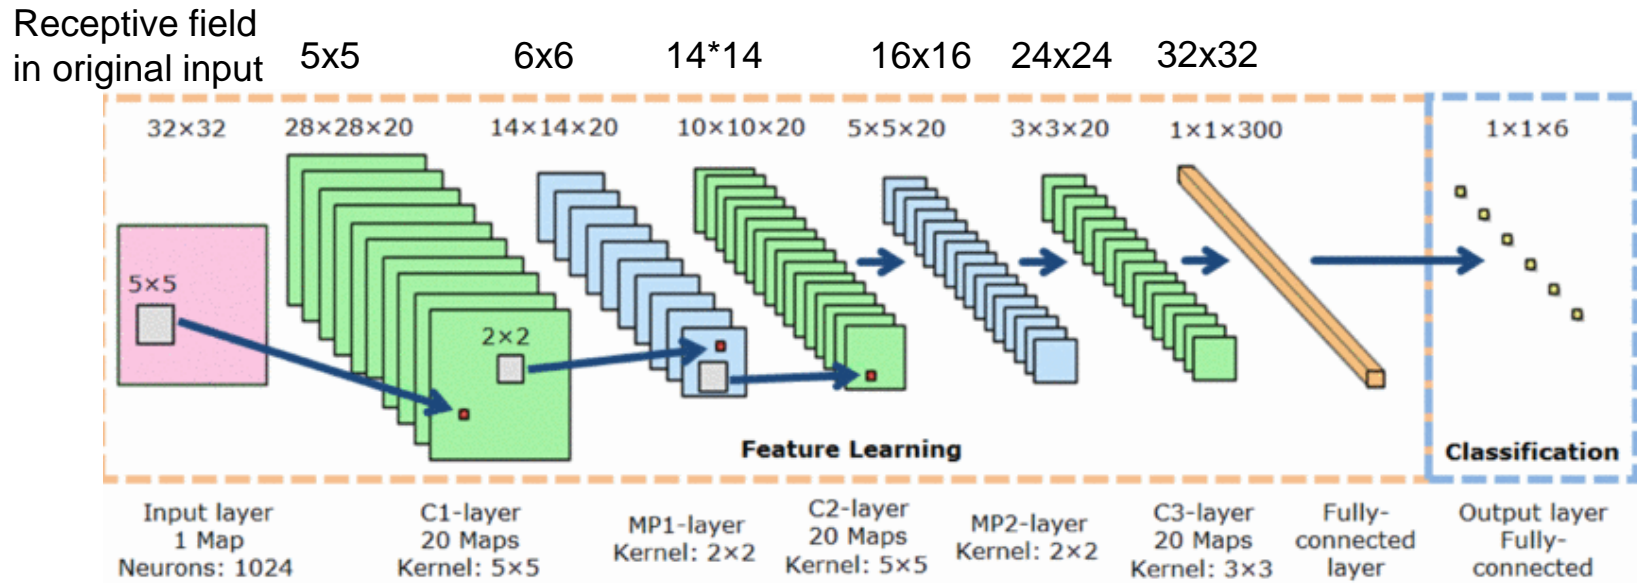




What input does activate a feature map in the CNN part or a neuron in the last layer?

# The receptive field is growing from layer to layer

The receptive field of a neuron is the area in the original input image that impact the value of this neuron – “that can be seen by this neuron”.




Neurons from feature maps in higher layers have a larger receptive field than neurons sitting in feature maps closer to the input.

Code to determine size of receptive field: <http://stackoverflow.com/questions/35582521/how-to-calculate-receptive-field-size>

# Visualize patches yielding high values in activation maps



**Figure 4: Top regions for six  $\text{pool}_5$  units.** Receptive fields and activation values are drawn in white. Some units are aligned to concepts, such as people (row 1) or text (4). Other units capture texture and material properties, such as dot arrays (2) and specular reflections (6).



<http://cs231n.github.io/understanding-cnn/>

Here we show image patches that activate maps in layer 5 most.

256 feature maps generated by 256 different 3x3 filters. Each feature map consists of equivalent neurons looking at different positions of the input volume.



# What kind of image (patches) excites a certain neuron corresponding to a large activation in a feature map?


10 images from data set leading to high signals 6 feature maps of **conv6**



10 images from data set leading to high signals 6 feature maps of **conv9**



# Homework: Develop a CNN for cifar10 data



Develop a CNN to classify cifar10 images (we have 10 classes)

Investigate the impact of standardizing the data on the performance

Notebook for homework:

[https://github.com/tensorchiefs/dl\\_course\\_2022/blob/master/notebooks/07\\_cifar10\\_norm.ipynb](https://github.com/tensorchiefs/dl_course_2022/blob/master/notebooks/07_cifar10_norm.ipynb)

# Summary

- Use loss curves to detect overfitting or underfitting problems
- NNs work best when respecting the underlying structure of the data.
  - Use fully connected NN for tabular data
  - Use convolutional NN for data with local order such as images
- CNNs exploit the local structure of images by local connections and shared weight (same kernel is applied at each position of the image).
- Use the relu activation function for hidden layers in CNNs.
- NNs are loosely inspired by the structure of the brain.
  - When going deep the receptive field increases (~layer 5 sees whole input)
  - Deeper layer respond to more complex feature in the input