

# Machine Intelligence:: Deep Learning

## Week 7

*Beate Sick, Loran Avci, Oliver Dürr*



Institut für Datenanalyse und Prozessdesign  
Zürcher Hochschule für Angewandte Wissenschaften

Part II: Bayesian NN via MC Dropout

Winterthur, 6. April. 2021

# Dropout








# Recall: Classical Dropout only during training





Using dropout during training implies:

- In each training step only weights to not-dropped units are updated  $\rightarrow$  we train a sparse sub-model NN
- For non-Bayesian NN we freeze the weights after training to a value  $w \cdot p^*$

# Recall: Dropout fights overfitting in a CIFAR10 CNN




# Recall: Nice properties of CNNs




**CNNs yield high accuracy and calibrated probabilities, but...**

# A non-Bayesian NN cannot ring the alarm

What happens if we present a novel class to the CNN?



**Plain wrong !**



**We need some error bars!**

From Dropout during training  
to MC Dropout during test time


# Bayesian NN via MC Dropout

Yarin Gal et al. (2015):



Via Dropout training we learned a whole weight distribution for each connection.

We can **sample from this Bernoulli-kind weight distribution** by performing **dropout during test time** and use the dropout-trained NN as Bayesian NN.

Gal showed that doing dropout approximates VI with a Bernoulli-kind variational distribution  $q_\theta$  (instead of a Gaussian).



Dropout in  
test time



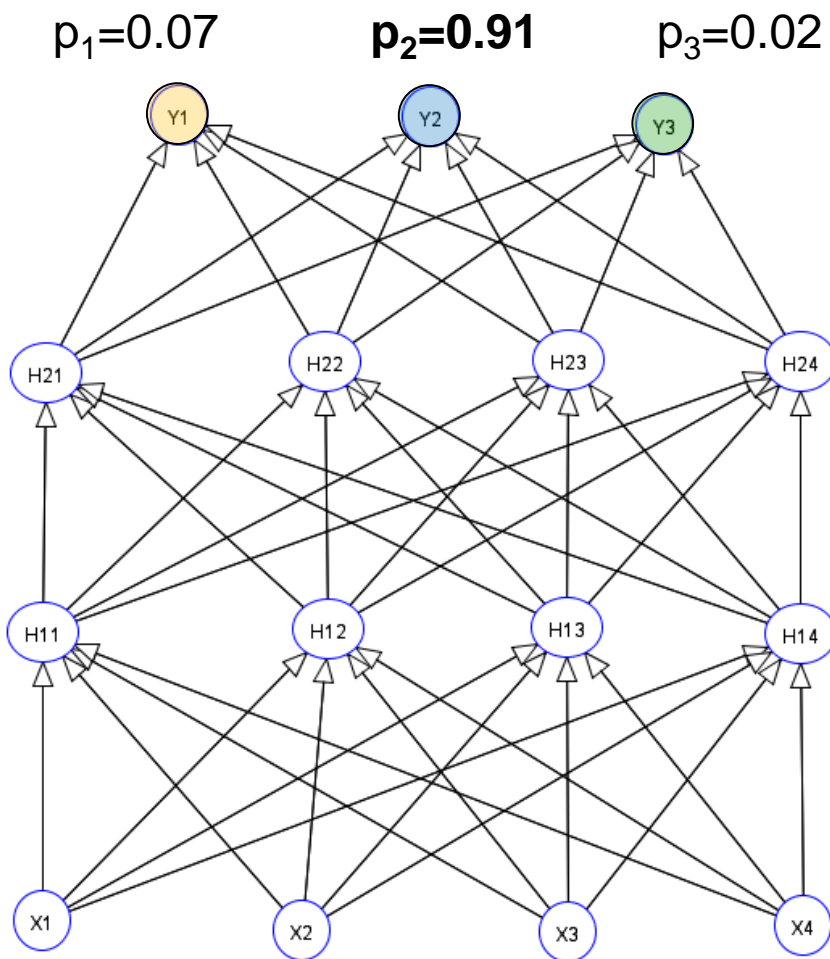
Which parameter has this  $q_\theta$ ?

The value  $w$ .



# When using Dropout only during training



For non-Bayesian NN we freeze the weights after training to a value  $w \cdot p^*$  and use then the trained NN for prediction:



Probability of predicted class:  $p_{\max}$


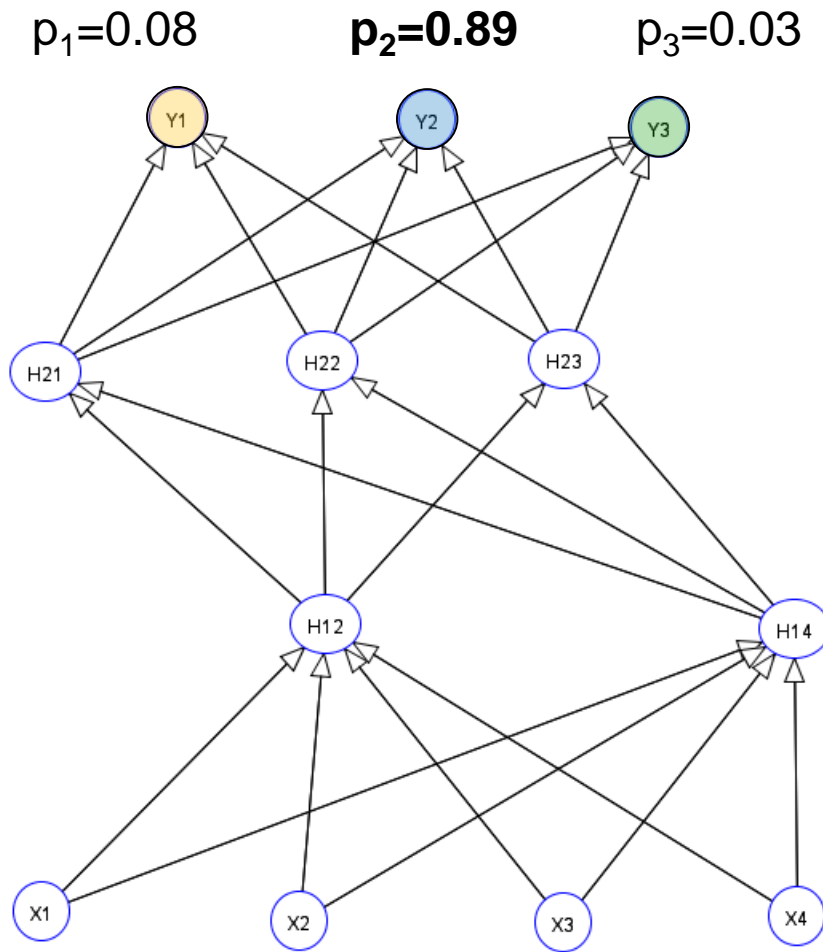
Input: image pixel values

# MC Dropout: we also perform dropout during test time



In each prediction instance we dropout a random subset of nodes, which corresponds to setting all weights starting from these nodes to zero.



# MC Dropout during test time: Run 1



Stochastic dropout of units

Same input image


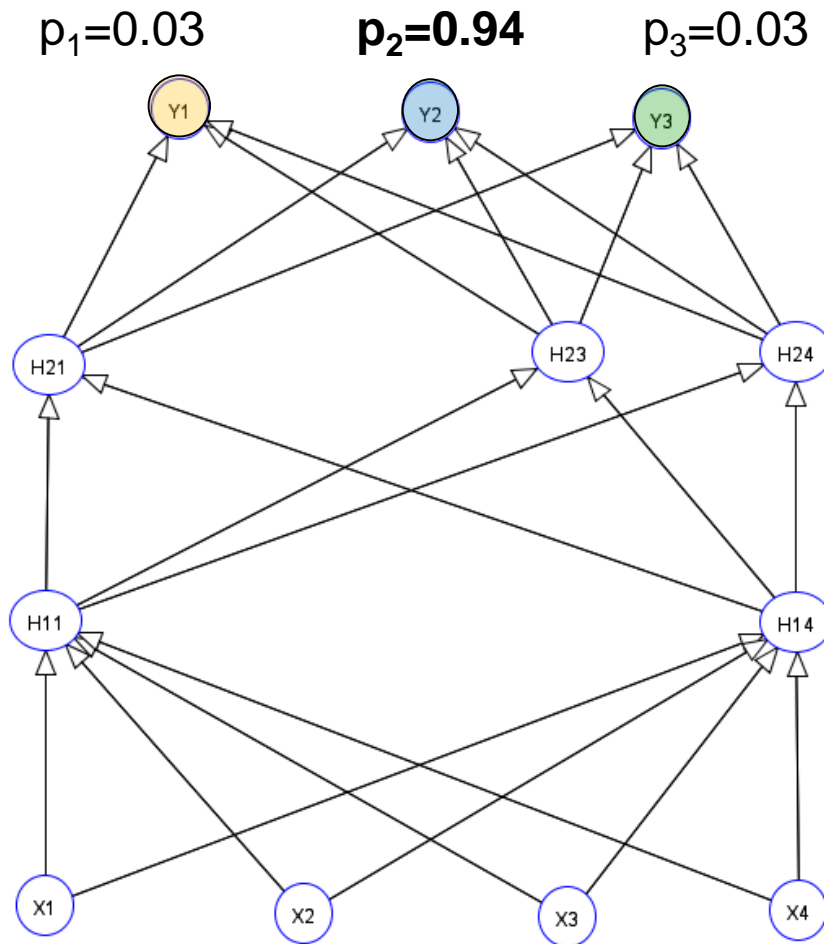
# MC Dropout during test time: Run 2



Stochastic dropout of units

Same input image


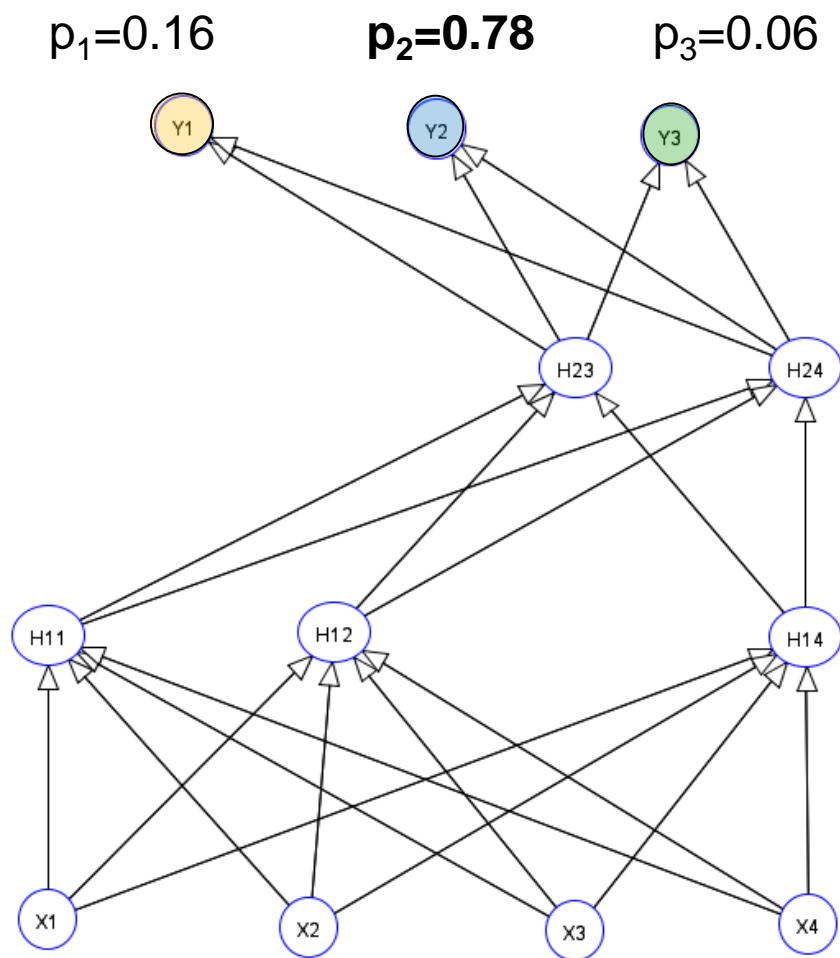
# MC Dropout during test time: Run 3



Stochastic dropout of units

Same input image

# MC Dropout during test time: Run 4



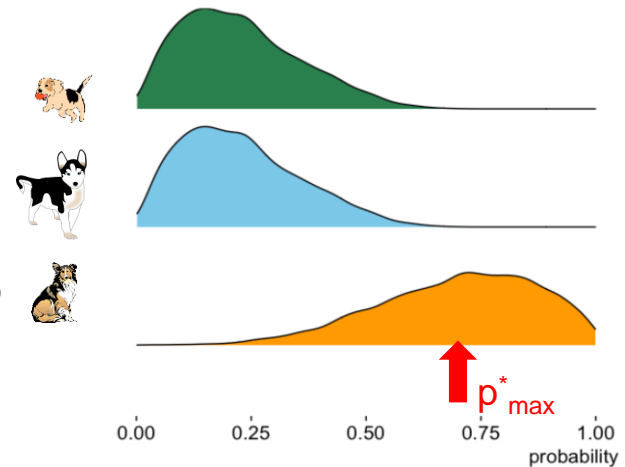

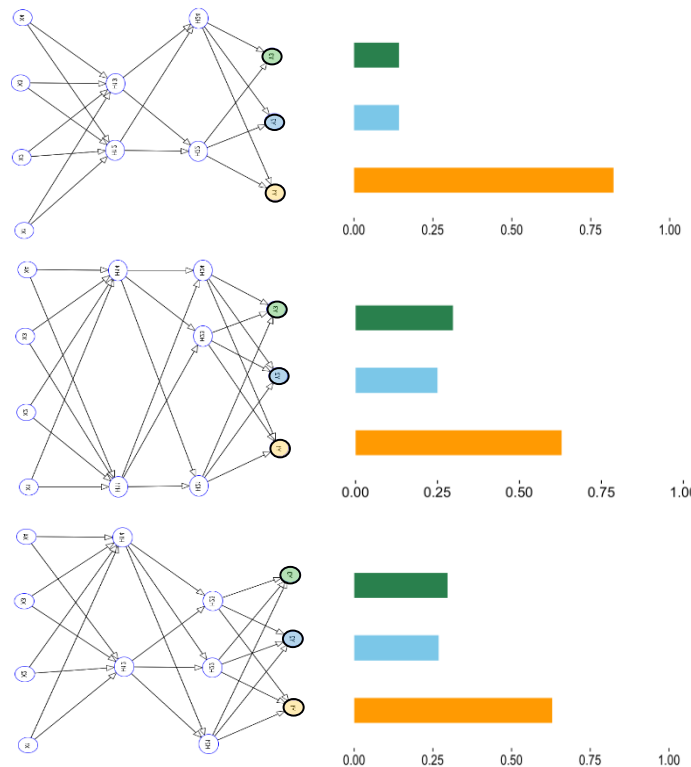
Stochastic dropout of units

Same input image

# MC Dropout during test time yields a multivariate predictive distribution for the parameters



Many Dropout Runs in forward pass




CNN predicts class “collie”  
but with high uncertainty

...

Remark: Mean of marginal give components of mean in multivariate distribution.

# Experiment with unknown phenotype




Dürr O, Murina E, Siegmund D, Tolkachev V, Steigle S, Sick B. Know when you don't know, Assay Drug Dev Technol. 2018



# Probability distribution from MC dropout runs

## Image with known class 15


100 MC predictions for an image with known phenotype 15



# Probability distribution from MC dropout runs


## Image with known class 15

100 MC predictions for an image with known phenotype 15



## Image with **unknown** class


100 MC predictions for an image with an unknown phenotype



# Comparing non-Bayesian with Bayesian NN


# Non-Bayesian and Bayesian NNs

Non-Bayesian  
NN




**Weights are fixed**

MC dropout  
Bayesian NN




**Weights have  
Bernoulli-kind  
distribution**

VI  
Bayesian NN



**Weights have  
Gaussian  
distribution**

# Comparing different Network types




A Non-Bayesian NN learns one set of weights: the same input same output  
A Bayesian NN learns distribution of weights: same input different outputs

# Uncertainty measures in classification

# Uncertainty in non-Bayesian classification

Multinomial CPD

$$MN(p_1(x, w), p_2(x, w), \dots, p_9(x, w))$$



In a non-Bayesian NN we make for each input  $x$  ONE CPD:

Image $x$
$MN(p_1(x, w), \dots, p_9(x, w))$

**Uncertainty** measures capturing the **aleatoric** uncertainty :

Negative log-Likelihood:  $NLL = -\log(p_{pred})$

Entropy:  $H = -\sum_{k=1}^9 p_k \cdot \log(p_k)$

# Uncertainty in Bayesian classification

In a Bayesian NN we sample T-times from the weight distributions and get each time a slightly different multinomial CPD

predict_no	Image x
1	MN(p1(x,w1), ..., p9(x,w1))
2	MN(p1(x,w2), ..., p9(x,w2))
...	
T	MN(p1(x,wT), ..., p9(x,wT))

For each class  $k$  ( $k \in \{1, 2, \dots, 9\}$ ) we determine the mean probability:  $p_k^* = \frac{1}{T} \sum_{i=1}^T p_{k_i}$

The predicted class has the highest mean probability:  $p_{pred}^* = \max(p_k^*)$


**Uncertainty** measures including **aleatoric** and **epistemic** contributions:

Entropy:  $H^* = - \sum_{k=1}^9 p_k^* \cdot \log(p_k^*)$

Total variance:  $V_{tot}^* = \sum_{k=1}^9 var(p_k) = \sum_{k=1}^9 \sum_{i=1}^T (p_{kt} - p_k^*)^2$



# Hands-on Time





Train a CNN with only 9 of the 10 classes and investigate if the uncertainties are different when predicting images from known or unknown classes.



[https://github.com/tensorchiefs/dl\\_course\\_2021/blob/master/notebooks/20\\_cifar10\\_classification\\_mc\\_and\\_vi.ipynb](https://github.com/tensorchiefs/dl_course_2021/blob/master/notebooks/20_cifar10_classification_mc_and_vi.ipynb)

# Looking at the predictive distribution!



Input image





Non-Bayesian CNN







Bayesian CNN via VI



Bayesian CNN via dropout



# Do known/novel classes yield different values for probability and uncertainty measures?







# Filtering experiment based on uncertainty

Goal: Get higher accuracy by filter only predictions which are quite certainly correct

- Each prediction has an attached uncertainty measure
- Sort predictions according to the uncertainty measures
- A set of predictions with very low uncertainties should achieve a high accuracy
- By successively adding predictions with increasing uncertainties should yield predictions sets with decreasing accuracies.

# Filtering experiment to compare uncertainty measures



Uncertainty from non-Bayesian NN is less good in filtering out wrong classifications than uncertainty measures from Bayesian variants of the NN.

# Uncertainty measures in regression

# Uncertainty in non-Bayesian NN

We do predictions for 400 x-values between -10 and 30 yielding for each x a Gaussian CPD.

x1= -10	x2= -9.9	...	x400= 30
$N(\mu_{x1,w}, \sigma_{x1,w})$	$N(\mu_{x2,w}, \sigma_{x2,w})$		$N(\mu_{x400,w}, \sigma_{x400,w})$

**Uncertainty** measures capturing the **aleatoric** uncertainty at  $x$ :

Standard deviation:  $\sigma_x$

95% PI:  $[q_{0.025}; q_{0.975}] = [\mu_x - 1.96 \cdot \sigma_x; \mu_x + 1.96 \cdot \sigma_x]$

Remark:


We could also estimate the 95% PI at position  $x$  by sampling several times from the CPD and determine the 0.025 and 0.975 quantiles, yielding :

95% PI:  $[q_{0.025}; q_{0.975}]$

# The problem of non-Bayesian NN

Problem:

A non-Bayesian NN does extrapolation with very small uncertainty





# Uncertainty in Bayesian regression NN

In a Bayesian NN we sample  $T$ -times from the weight distributions and get each time a slightly different CPD. In regression the CPD is often Gaussian.

We do predictions for 400  $x$ -values between -10 and 30 yielding in each of the  $T$  runs a different Gaussian CPD at each  $x$ -position.

<b>predict_no</b>	$x_1 = -10$	$x_2 = -9.9$	...	$x_{400} = 30$
1	$N(x_1, w_1, x_1, w_1)$	$N(x_2, w_1, x_2, w_1)$		$N(x_{400}, w_1, x_{400}, w_1)$
2	$N(x_1, w_2, x_1, w_2)$	$N(x_2, w_2, x_2, w_2)$		$N(x_{400}, w_2, x_{400}, w_2)$
...				
$T$	$N(x_1, w_T, x_1, w_T)$	$N(x_2, w_T, x_2, w_T)$		$N(x_{400}, w_T, x_{400}, w_T)$


**Uncertainty** measures including **aleatoric** and **epistemic** contributions:

To estimate the 95% PI at position  $x$  we sample  $y$ -values from each of the  $T$  CPDs and determine from the samples the 0.025 and 0.975 quantiles, yielding :


95% PI:  $[q_{0.025}; q_{0.975}]$

# Can we see enhanced uncertainty in extrapolation


The solid lines show five predicted y-vectors corresponding to 5 CPDs at each x-position.






<https://youtu.be/FO5avm3XT4g>



<https://youtu.be/mQrUcUoT2k4>



<https://youtu.be/0-oyDeR9HrE>



# Conclusion

- Standard neural networks (NNs) fail to express their uncertainty (can't talk about the elephant in the room).
- Bayesian neural networks (BNNs) can express their uncertainty.
- BNNs often yield better performance than their non-Bayesian variant.
- Novel classes can be better identified with BNNs, which combine epistemic and aleatoric uncertainties compared to standard NNs.
- Variational inference (VI) and Monte Carlo dropout (MC dropout) are approximation methods that allow you to fit deep BNNs.
- TFP provides easy to use layers for fitting a BNN via VI.
- MC dropout can be used in Keras for fitting BNNs.